

Table of Contents

[Overview](#)

[New and updated articles](#)

[Data types](#)

[Database Console Commands](#)

[Functions](#)

[Language elements](#)

[Queries](#)

[Statements](#)

[Tutorial: Writing Transact-SQL Statements](#)

[Lesson 1: Creating Database Objects](#)

[Lesson 1-1 - Creating a Database](#)

[Lesson 1-2 - Creating a Table](#)

[Lesson 1-3 - Inserting and Updating Data in a Table](#)

[Lesson 1-4 - Reading the Data in a Table](#)

[Lesson 1-5 - Summary - Creating Database Objects](#)

[Lesson 2: Configuring Permissions on Database Objects](#)

[Lesson 2-1 - Creating a Login](#)

[Lesson 2-2 - Granting Access to a Database](#)

[Lesson 2-3 - Creating Views and Stored Procedures](#)

[Lesson 2-4 - Granting Access to a Database Object](#)





[Lesson 2-5 - Summary - Configuring Permissions on Database Objects](#)

[Lesson 3: Deleting Database Objects](#)

[Lesson 3-1 - Deleting Database Objects](#)

Transact-SQL Reference (Database Engine)

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

This topic gives the basics about how to find and use the Microsoft Transact-SQL (T-SQL) reference topics. T-SQL is central to using Microsoft SQL products and services. All tools and applications that communicate with a SQL database do so by sending T-SQL commands.

Tools that use T-SQL

Some of the Microsoft tools that issue T-SQL commands are:

- [SQL Server Management Studio](#)
- [SQL Server Data Tools](#).
- [sqlcmd](#).
- [SQL Operations Studio \(preview\)](#).

Locate the Transact-SQL reference topics

To find T-SQL topics, use search at the top right of this page, or use the table of contents on the left side of the page. You can also type a T-SQL key word in the Management Studio Query Editor window, and press F1.

Find system views





To find the system tables, views, functions, and procedures, see these links which are in the [Using relational databases](#) section of the SQL documentation.

- [System catalog Views](#)
- [System compatibility views](#)
- [System dynamic management views](#)
- [System functions](#)
- [System information schema views](#)
- [System stored procedures](#)
- [System tables](#)

"Applies to" references

The T-SQL reference topics encompass multiple versions of SQL Server, starting with 2008, as well as the other Azure SQL services. Near the top of each topic is a section that indicates which products and services support subject of the topic.

For example, this topic applies to all versions, and has the following label.

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Another example, the following label indicates a topic that applies only to Azure SQL Data Warehouse and Parallel Data Warehouse.

THIS TOPIC APPLIES TO:  SQL Server  Azure SQL Database  Azure SQL Data Warehouse  Parallel

In some cases, the topic is used by a product or service, but all of the arguments are not supported. In this case, additional **Applies to** sections are inserted into the appropriate argument descriptions in the body of the topic.

Get help from the MSDN forum

For online help, see the [MSDN Transact-SQL Forum](#).

See other language references

The SQL docs include these other language references:

- [XQuery Language Reference](#)
- [Integration Services Language Reference](#)
- [Replication Language Reference](#)
- [Analysis Services Language Reference](#)

Next steps

Now that you understand how to find the T-SQL reference topics, you are ready to:

- Work through a short tutorial about how to write T-SQL, see [Tutorial: Writing Transact-SQL Statements](#).
- View the [Transact-SQL Syntax Conventions \(Transact-SQL\)](#).

New and Recently Updated: Transact-SQL docs

5/1/2018 • 4 min to read • [Edit Online](#)

Nearly every day Microsoft updates some of its existing articles on its [Docs.Microsoft.com](#) documentation website. This article displays excerpts from recently updated articles. Links to new articles might also be listed.

This article is generated by a program that is rerun periodically. Occasionally an excerpt can appear with imperfect formatting, or as markdown from the source article. Images are never displayed here.

Recent updates are reported for the following date range and subject:

- *Date range of updates:* **2018-02-03** -to- **2018-04-28**
- *Subject area:* **T-SQL**.

New Articles Created Recently

The following links jump to new articles that have been added recently.

There are no new articles to list, this time.

Updated Articles with Excerpts

This section displays the excerpts of updates gathered from articles that have recently experienced a large update.

The excerpts displayed here appear separated from their proper semantic context. Also, sometimes an excerpt is separated from important markdown syntax that surrounds it in the actual article. Therefore these excerpts are for general guidance only. The excerpts only enable you to know whether your interests warrant taking the time to click and visit the actual article.

For these and other reasons, do not copy code from these excerpts, and do not take as exact truth any text excerpt. Instead, visit the actual article.

Compact List of Articles Updated Recently

This compact list provides links to all the updated articles that are listed in the Excerpts section.

1. [ALTER DATABASE SCOPED CONFIGURATION \(Transact-SQL\)](#)
2. [RESTORE Statements \(Transact-SQL\)](#)

1. [ALTER DATABASE SCOPED CONFIGURATION \(Transact-SQL\)](#)

Updated: 2018-04-13 ([Next](#))

XTP_PROCEDURE_EXECUTION_STATISTICS = { ON | **OFF** }

Applies to: {Included-Content-Goes-Here}

Enables or disables collection of execution statistics at the module-level for natively compiled T-SQL modules in the current database. The default is OFF. The execution statistics are reflected in [sys.dm_exec_procedure_stats].

Module-level execution statistics for natively compiled T-SQL modules are collected if either this option is ON, or if statistics collection is enabled through [sp_xtp_control_proc_exec_stats].

XTP_QUERY_EXECUTION_STATISTICS = { ON | **OFF** }

Applies to: {Included-Content-Goes-Here}

Enables or disables collection of execution statistics at the statement-level for natively compiled T-SQL modules in the current database. The default is OFF. The execution statistics are reflected in [sys.dm_exec_query_stats] and in [Query Store].

Statement-level execution statistics for natively compiled T-SQL modules are collected if either this option is ON, or if statistics collection is enabled through [sp_xtp_control_query_exec_stats].

For more details about performance monitoring of natively-compiled T-SQL modules see [Monitoring Performance of Natively Compiled Stored Procedures].

2. RESTORE Statements (Transact-SQL)

Updated: 2018-04-13 ([Previous](#))

General Remarks - SQL Database Managed Instance

For an asynchronous restore, the restore continues even if client connection breaks. If your connection is dropped, you can check [sys.dm_operation_status] view for the status of a restore operation (as well as for CREATE and DROP database).

The following database options are set/overridden and cannot be changed later:

- NEW_BROKER (if broker is not enabled in .bak file)
- ENABLE_BROKER (if broker is not enabled in .bak file)
- AUTO_CLOSE=OFF (if a database in .bak file has AUTO_CLOSE=ON)
- RECOVERY FULL (if a database in .bak file has SIMPLE or BULK_LOGGED recovery mode)
- Memory optimized filegroup is added and called XTP if it was not in the source .bak file. Any existing memory optimized filegroup is renamed to XTP
- SINGLE_USER and RESTRICTED_USER options are converted to MULTI_USER

Limitations - SQL Database Managed Instance

These limitations apply:

- .BAK files containing multiple backup sets cannot be restored.
- .BAK files containing multiple log files cannot be restored.
- Restore will fail if .bak contains FILESTREAM data.
- Backups containing databases that have active In-memory objects cannot currently be restored.
- Backups containing databases where at some point In-Memory objects existed cannot currently be restored.
- Backups containing databases in read-only mode cannot currently be restored. This limitation will be removed soon.

For more information, see [Managed Instance]

Similar articles about new or updated articles

This section lists very similar articles for recently updated articles in other subject areas, within our public GitHub.com repository: [MicrosoftDocs/sql-docs](#).

Subject areas that *do* have new or recently updated articles





- New + Updated (11+6): [Advanced Analytics for SQL](#) docs
- New + Updated (18+0): [Analysis Services for SQL](#) docs
- New + Updated (218+14): [Connect to SQL](#) docs
- New + Updated (14+0): [Database Engine for SQL](#) docs
- New + Updated (3+2): [Integration Services for SQL](#) docs
- New + Updated (3+3): [Linux for SQL](#) docs
- New + Updated (7+10): [Relational Databases for SQL](#) docs
- New + Updated (0+2): [Reporting Services for SQL](#) docs
- New + Updated (1+3): [SQL Operations Studio](#) docs
- New + Updated (2+3): [Microsoft SQL Server](#) docs
- New + Updated (1+1): [SQL Server Data Tools \(SSDT\)](#) docs
- New + Updated (5+2): [SQL Server Management Studio \(SSMS\)](#) docs
- New + Updated (0+2): [Transact-SQL](#) docs
- New + Updated (1+1): [Tools for SQL](#) docs

Subject areas that do *not* have any new or recently updated articles

- New + Updated (0+0): [Analytics Platform System for SQL](#) docs
- New + Updated (0+0): [Data Quality Services for SQL](#) docs
- New + Updated (0+0): [Data Mining Extensions \(DMX\) for SQL](#) docs
- New + Updated (0+0): [Master Data Services \(MDS\) for SQL](#) docs
- New + Updated (0+0): [Multidimensional Expressions \(MDX\) for SQL](#) docs
- New + Updated (0+0): [ODBC \(Open Database Connectivity\) for SQL](#) docs
- New + Updated (0+0): [PowerShell for SQL](#) docs
- New + Updated (0+0): [Samples for SQL](#) docs
- New + Updated (0+0): [SQL Server Migration Assistant \(SSMA\)](#) docs
- New + Updated (0+0): [XQuery for SQL](#) docs

Data types (Transact-SQL)

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

In SQL Server, each column, local variable, expression, and parameter has a related data type. A data type is an attribute that specifies the type of data that the object can hold: integer data, character data, monetary data, date and time data, binary strings, and so on.

SQL Server supplies a set of system data types that define all the types of data that can be used with SQL Server. You can also define your own data types in Transact-SQL or the Microsoft .NET Framework. Alias data types are based on the system-supplied data types. For more information about alias data types, see [CREATE TYPE \(Transact-SQL\)](#). User-defined types obtain their characteristics from the methods and operators of a class that you create by using one of the programming languages support by the .NET Framework.

When two expressions that have different data types, collations, precision, scale, or length are combined by an operator, the characteristics of result are determined by the following:

- The data type of the result is determined by applying the rules of data type precedence to the data types of the input expressions. For more information, see [Data Type Precedence \(Transact-SQL\)](#).
- The collation of the result is determined by the rules of collation precedence when the result data type is **char**, **varchar**, **text**, **nchar**, **nvarchar**, or **ntext**. For more information, see [Collation Precedence \(Transact-SQL\)](#).
- The precision, scale, and length of the result depend on the precision, scale, and length of the input expressions. For more information, see [Precision, Scale, and Length \(Transact-SQL\)](#).

SQL Server provides data type synonyms for ISO compatibility. For more information, see [Data Type Synonyms \(Transact-SQL\)](#).

Data type categories

Data types in SQL Server are organized into the following categories:

Exact numerics	Unicode character strings
Approximate numerics	Binary strings
Date and time	Other data types
Character strings	

In SQL Server, based on their storage characteristics, some data types are designated as belonging to the following groups:

- Large value data types: **varchar(max)**, and **nvarchar(max)**
- Large object data types: **text**, **ntext**, **image**, **varbinary(max)**, and **xml**

NOTE

sp_help returns -1 as the length for the large-value and **xml** data types.

Exact numerics

bigint	numeric
bit	smallint
decimal	smallmoney
int	tinyint
money	

Approximate numerics

float	real

Date and time

date	datetimeoffset
datetime2	smalldatetime
datetime	time

Character strings

char	varchar
text	

Unicode character strings

nchar	nvarchar
ntext	

Binary strings

binary	varbinary
image	

Other data types

cursor	rowversion

hierarchyid	uniqueidentifier
sql_variant	xml
Spatial Geometry Types	Spatial Geography Types
table	

See also

[CREATE PROCEDURE \(Transact-SQL\)](#)

[CREATE TABLE \(Transact-SQL\)](#)

[DECLARE @local_variable \(Transact-SQL\) EXECUTE \(Transact-SQL\)](#)

[Expressions \(Transact-SQL\)](#)

[Functions \(Transact-SQL\)](#)

[LIKE \(Transact-SQL\)](#)





[sp_droptype \(Transact-SQL\)](#)

[sp_help \(Transact-SQL\)](#)

[sp_rename \(Transact-SQL\)](#)

Database Console Commands

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse





SQL Server provides the following management commands.

In This Section

CHECKPOINT	KILL STATS JOB
DBCC	RECONFIGURE
KILL	SHUTDOWN
KILL QUERY NOTIFICATION SUBSCRIPTION	

What are the SQL database functions?

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Learn about the categories of built-in functions you can use with SQL databases. You can use the built-in functions or create your own user-defined functions.

Aggregate functions

Aggregate functions perform a calculation on a set of values and return a single value. They are allowed in the select list or the HAVING clause of a SELECT statement. You can use an aggregation in combination with the GROUP BY clause to calculate the aggregation on categories of rows. Use the OVER clause to calculate the aggregation on a specific range of value. The OVER clause cannot follow the GROUPING or GROUPING_ID aggregations.

All aggregate functions are deterministic, which means they always return the same value when they run on the same input values. For more information, see [Deterministic and Nondeterministic Functions](#).

Analytic functions

Analytic functions compute an aggregate value based on a group of rows. However, unlike aggregate functions, analytic functions can return multiple rows for each group. You can use analytic functions to compute moving averages, running totals, percentages, or top-N results within a group.

Ranking functions

Ranking functions return a ranking value for each row in a partition. Depending on the function that is used, some rows might receive the same value as other rows. Ranking functions are nondeterministic.

Rowset functions

Rowset functions Return an object that can be used like table references in an SQL statement.

Scalar functions

Operate on a single value and then return a single value. Scalar functions can be used wherever an expression is valid.

Categories of scalar functions

FUNCTION CATEGORY	DESCRIPTION
Configuration Functions	Return information about the current configuration.
Conversion Functions	Support data type casting and converting.
Cursor Functions	Return information about cursors.

FUNCTION CATEGORY	DESCRIPTION
Date and Time Data Types and Functions	Perform operations on a date and time input values and return string, numeric, or date and time values.
JSON Functions	Validate, query, or change JSON data.
Logical Functions	Perform logical operations.
Mathematical Functions	Perform calculations based on input values provided as parameters to the functions, and return numeric values.
Metadata Functions	Return information about the database and database objects.
Security Functions	Return information about users and roles.
String Functions	Perform operations on a string (char or varchar) input value and return a string or numeric value.
System Functions	Perform operations and return information about values, objects, and settings in an instance of SQL Server.
System Statistical Functions	Return statistical information about the system.
Text and Image Functions	Perform operations on text or image input values or columns, and return information about the value.

Function Determinism

SQL Server built-in functions are either deterministic or nondeterministic. Functions are deterministic when they always return the same result any time they are called by using a specific set of input values. Functions are nondeterministic when they could return different results every time they are called, even with the same specific set of input values. For more information, see [Deterministic and Nondeterministic Functions](#)

Function Collation

Functions that take a character string input and return a character string output use the collation of the input string for the output.

Functions that take non-character inputs and return a character string use the default collation of the current database for the output.

Functions that take multiple character string inputs and return a character string use the rules of collation precedence to set the collation of the output string. For more information, see [Collation Precedence \(Transact-SQL\)](#).

See Also





[CREATE FUNCTION \(Transact-SQL\)](#)

[Deterministic and Nondeterministic Functions](#)

[Using Stored Procedures \(MDX\)](#)

Language Elements (Transact-SQL)

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

SQL Server supports the following language elements.

In This Section

[-- \(Comment\) \(Transact-SQL\)](#)

[Slash Star \(Block Comment\) \(Transact-SQL\)](#)

[CREATE DIAGNOSTICS SESSION \(Transact-SQL\)](#)





[NULL and UNKNOWN \(Transact-SQL\)](#)

[Transactions \(SQL Data Warehouse\)](#)

[USE \(Transact-SQL\)](#)

Queries

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Data Manipulation Language (DML) is a vocabulary used to retrieve and work with data in SQL Server 2017 and SQL Database. Most also work in SQL Data Warehouse and PDW (review each individual statement for details). Use these statements to add, modify, query, or remove data from a SQL Server database.

In This Section

The following table lists the DML statements that SQL Server uses.





BULK INSERT (Transact-SQL)	SELECT (Transact-SQL)
DELETE (Transact-SQL)	UPDATE (Transact-SQL)
INSERT (Transact-SQL)	UPDATETEXT (Transact-SQL)
MERGE (Transact-SQL)	WRITETEXT (Transact-SQL)
READTEXT (Transact-SQL)	

The following table lists the clauses that are used in multiple DML statements or clauses.

CLAUSE	CAN BE USED IN THESE STATEMENTS
FROM (Transact-SQL)	DELETE, SELECT, UPDATE
Hints (Transact-SQL)	DELETE, INSERT, SELECT, UPDATE
OPTION Clause (Transact-SQL)	DELETE, SELECT, UPDATE
OUTPUT Clause (Transact-SQL)	DELETE, INSERT, MERGE, UPDATE
Search Condition (Transact-SQL)	DELETE, MERGE, SELECT, UPDATE
Table Value Constructor (Transact-SQL)	FROM, INSERT, MERGE
TOP (Transact-SQL)	DELETE, INSERT, MERGE, SELECT, UPDATE
WHERE (Transact-SQL)	DELETE, SELECT, UPDATE, MATCH
WITH common_table_expression (Transact-SQL)	DELETE, INSERT, MERGE, SELECT, UPDATE

Transact-SQL statements

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

This reference topic summarizes the categories of statements for use with Transact-SQL (T-SQL). You can find all of the statements listed in the left-hand navigation.

Backup and restore

The backup and restore statements provide ways to create backups and restore from backups. For more information, see the [Backup and restore overview](#).

Data Definition Language

Data Definition Language (DDL) statements defines data structures. Use these statements to create, alter, or drop data structures in a database.

- ALTER
- Collations
- CREATE
- DROP
- DISABLE TRIGGER
- ENABLE TRIGGER
- RENAME
- UPDATE STATISTICS

Data Manipulation Language

Data Manipulation Language (DML) affect the information stored in the database. Use these statements to insert, update, and change the rows in the database.

- BULK INSERT
- DELETE
- INSERT
- MERGE
- TRUNCATE TABLE

Permissions statements

Permissions statements determine which users and logins can access data and perform operations. For more information about authentication and access, see the [Security center](#).

Service Broker statements





Service Broker is a feature that provides native support for messaging and queuing applications. For more information, see [Service Broker](#).

Session settings

SET statements determine how the current session handles run time settings. For an overview, see [SET statements](#).

Tutorial: Writing Transact-SQL Statements

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Welcome to the Writing Transact-SQL Statements tutorial. This tutorial is intended for users who are new to writing SQL statements. It will help new users get started by reviewing some basic statements for creating tables and inserting data. This tutorial uses Transact-SQL, the Microsoft implementation of the SQL standard. This tutorial is intended as a brief introduction to the Transact-SQL language and not as a replacement for a Transact-SQL class. The statements in this tutorial are intentionally simple, and are not meant to represent the complexity found in a typical production database.

NOTE: If you are a beginner you might find it easier to use SQL Server Management Studio instead of writing Transact-SQL statements.

Finding More Information

To find more information about any specific statement, either search for the statement by name in SQL Server Books Online, or use the Contents to browse the 1,800 language elements listed alphabetically under [Transact-SQL Reference \(Database Engine\)](#). Another good strategy for finding information is to search for key words that are related to the subject matter you are interested in. For example, if you want to know how to return a part of a date (such as the month), search the index for **dates [SQL Server]**, and then select **dateparts**. This takes you to the topic [DATEPART \(Transact-SQL\)](#). As another example, to find out how to work with strings, search for **string functions**. This takes you to the topic [String Functions \(Transact-SQL\)](#).

What You Will Learn

This tutorial shows you how to create a database, create a table in the database, insert data into the table, update the data, read the data, delete the data, and then delete the table. You will create views and stored procedures and configure a user to the database and the data.

This tutorial is divided into three lessons:

[Lesson 1: Creating Database Objects](#)

In this lesson, you create a database, create a table in the database, insert data into the table, update the data, and read the data.

[Lesson 2: Configuring Permissions on Database Objects](#)

In this lesson, you create a login and user. You will also create a view and a stored procedure, and then grant the user permission to the stored procedure.

[Lesson 3: Deleting Database Objects](#)

In this lesson, you remove access to data, delete data from a table, delete the table, and then delete the database.

Requirements





To complete this tutorial, you do not have to know the SQL language, but you should understand basic database concepts such as tables. During this tutorial, you will create a database and create a Windows user. These tasks require a high level of permissions; therefore, you should log in to the computer as an administrator.

Your system must have the following installed:

- Any edition of SQL Server.
- [SQL Server Management Studio](#)

Lesson 1: Creating Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

This lesson shows you how to create a database, create a table in the database, and then access and change the data in the table. Because this lesson is an introduction to using Transact-SQL, it does not use or describe the many options that are available for these statements.

Transact-SQL statements can be written and submitted to the Database Engine in the following ways:

- By using SQL Server Management Studio. This tutorial assumes that you are using Management Studio, but you can also use Management Studio Express, which is available as a free download from the [Microsoft Download Center](#).
- By using the [sqlcmd utility](#).
- By connecting from an application that you create.

The code executes on the Database Engine in the same way and with the same permissions, regardless of how you submit the code statements.

To run Transact-SQL statements in Management Studio, open Management Studio and connect to an instance of the SQL Server Database Engine.

This lesson contains the following topics:





- [Creating a Database \(Tutorial\)](#)
- [Creating a Table \(Tutorial\)](#)
- [Inserting and Updating Data in a Table \(Tutorial\)](#)
- [Reading the Data in a Table \(Tutorial\)](#)
- [Summary: Creating Database Objects](#)

Next Task in Lesson

[Creating a Database \(Tutorial\)](#)

Lesson 1-1 - Creating a Database

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Like many Transact-SQL statements, the CREATE DATABASE statement has a required parameter: the name of the database. CREATE DATABASE also has many optional parameters, such as the disk location where you want to put the database files. When you execute CREATE DATABASE without the optional parameters, SQL Server uses default values for many of these parameters. This tutorial uses very few of the optional syntax parameters.

To create a database

1. In a Query Editor window, type but do not execute the following code:

```
CREATE DATABASE TestData  
GO
```

2. Use the pointer to select the words `CREATE DATABASE`, and then press **F1**. The CREATE DATABASE topic in SQL Server Books Online should open. You can use this technique to find the complete syntax for CREATE DATABASE and for the other statements that are used in this tutorial.
3. In Query Editor, press **F5** to execute the statement and create a database named `TestData`.

When you create a database, SQL Server makes a copy of the **model** database, and renames the copy to the database name. This operation should only take several seconds, unless you specify a large initial size of the database as an optional parameter.

NOTE

The keyword GO separates statements when more than one statement is submitted in a single batch. GO is optional when the batch contains only one statement.

Next Task in Lesson





[Creating a Table \(Tutorial\)](#)

See Also

[CREATE DATABASE \(SQL Server Transact-SQL\)](#)

Lesson 1-2 - Creating a Table

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

To create a table, you must provide a name for the table, and the names and data types of each column in the table. It is also a good practice to indicate whether null values are allowed in each column. To create a table, you must have the `CREATE TABLE` permission, and the `ALTER SCHEMA` permission on the schema that will contain the table. The `db_ddladmin` fixed database role has these permissions.

Most tables have a primary key, made up of one or more columns of the table. A primary key is always unique. The Database Engine will enforce the restriction that any primary key value cannot be repeated in the table.

For a list of data types and links for a description of each, see [Data Types \(Transact-SQL\)](#).

NOTE

The Database Engine can be installed as case sensitive or non-case sensitive. If the Database Engine is installed as case sensitive, object names must always have the same case. For example, a table named OrderData is a different table from a table named ORDERDATA. If the Database Engine is installed as non-case sensitive, those two table names are considered to be the same table, and that name can only be used one time.

To create a database to contain the new table

- Enter the following code into a Query Editor window.

```
USE master;
GO

--Delete the TestData database if it exists.
IF EXISTS(SELECT * from sys.databases WHERE name='TestData')
BEGIN
    DROP DATABASE TestData;
END

--Create a new database called TestData.
CREATE DATABASE TestData;
Press the F5 key to execute the code and create the database.
```

Switch the Query Editor connection to the TestData database

- In a Query Editor window, type and execute the following code to change your connection to the `TestData` database.

```
USE TestData
GO
```

To create a table

- In a Query Editor window, type and execute the following code to create a simple table named `Products`. The columns in the table are named `ProductID`, `ProductName`, `Price`, and `ProductDescription`. The `ProductID` column is the primary key of the table. `int`, `varchar(25)`, `money`, and `text` are all data types. Only the `Price` and `ProductDescription` columns can have no data when a row is inserted or changed.

This statement contains an optional element (`dbo.`) called a schema. The schema is the database object that owns the table. If you are an administrator, `dbo` is the default schema. `dbo` stands for database owner.

```
CREATE TABLE dbo.Products
(
    ProductID int PRIMARY KEY NOT NULL,
    ProductName varchar(25) NOT NULL,
    Price money NULL,
    ProductDescription text NULL)
GO
```

Next Task in Lesson





[Inserting and Updating Data in a Table \(Tutorial\)](#)

See Also

[CREATE TABLE \(Transact-SQL\)](#)

Lesson 1-3 - Inserting and Updating Data in a Table

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Now that you have created the **Products** table, you are ready to insert data into the table by using the INSERT statement. After the data is inserted, you will change the content of a row by using an UPDATE statement. You will use the WHERE clause of the UPDATE statement to restrict the update to a single row. The four statements will enter the following data.

PRODUCTID	PRODUCTNAME	PRICE	PRODUCTDESCRIPTION
1	Clamp	12.48	Workbench clamp
50	Screwdriver	3.17	Flat head
75	Tire Bar		Tool for changing tires.
3000	3mm Bracket	.52	

The basic syntax is: INSERT, table name, column list, VALUES, and then a list of the values to be inserted. The two hyphens in front of a line indicate that the line is a comment and the text will be ignored by the compiler. In this case, the comment describes a permissible variation of the syntax.

To insert data into a table

1. Execute the following statement to insert a row into the `Products` table that was created in the previous task. This is the basic syntax.

```
-- Standard syntax
INSERT dbo.Products (ProductID, ProductName, Price, ProductDescription)
VALUES (1, 'Clamp', 12.48, 'Workbench clamp')
GO
```

2. The following statement shows how you can change the order in which the parameters are provided by switching the placement of the `ProductID` and `ProductName` in both the field list (in parentheses) and in the values list.

```
-- Changing the order of the columns
INSERT dbo.Products (ProductName, ProductID, Price, ProductDescription)
VALUES ('Screwdriver', 50, 3.17, 'Flat head')
GO
```

3. The following statement demonstrates that the names of the columns are optional, as long as the values are listed in the correct order. This syntax is common but is not recommended because it might be harder for others to understand your code. `NULL` is specified for the `Price` column because the price for this product is not yet known.

```
-- Skipping the column list, but keeping the values in order
INSERT dbo.Products
    VALUES (75, 'Tire Bar', NULL, 'Tool for changing tires.')
GO
```

4. The schema name is optional as long as you are accessing and changing a table in your default schema.

Because the `ProductDescription` column allows null values and no value is being provided, the

`ProductDescription` column name and value can be dropped from the statement completely.

```
-- Dropping the optional dbo and dropping the ProductDescription column
INSERT Products (ProductID, ProductName, Price)
    VALUES (3000, '3mm Bracket', .52)
GO
```

To update the products table

1. Type and execute the following `UPDATE` statement to change the `ProductName` of the second product from `Screwdriver` to `Flat Head Screwdriver`.

```
UPDATE dbo.Products
    SET ProductName = 'Flat Head Screwdriver'
    WHERE ProductID = 50
GO
```

Next Task in Lesson

[Reading the Data in a Table \(Tutorial\)](#)





See Also

[INSERT \(Transact-SQL\)](#)

[UPDATE \(Transact-SQL\)](#)

Lesson 1-4 - Reading the Data in a Table

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Use the SELECT statement to read the data in a table. The SELECT statement is one of the most important Transact-SQL statements, and there are many variations in the syntax. For this tutorial, you will work with five simple versions.

To read the data in a table

1. Type and execute the following statements to read the data in the `Products` table.

```
-- The basic syntax for reading data from a single table
SELECT ProductID, ProductName, Price, ProductDescription
FROM dbo.Products
GO
```

2. You can use an asterisk to select all the columns in the table. This is often used in ad hoc queries. You should provide the column list in your permanent code so that the statement will return the predicted columns, even if a new column is added to the table later.

```
-- Returns all columns in the table
-- Does not use the optional schema, dbo
SELECT * FROM Products
GO
```

3. You can omit columns that you do not want to return. The columns will be returned in the order that they are listed.

```
-- Returns only two of the columns from the table
SELECT ProductName, Price
FROM dbo.Products
GO
```

4. Use a `WHERE` clause to limit the rows that are returned to the user.

```
-- Returns only two of the records in the table
SELECT ProductID, ProductName, Price, ProductDescription
FROM dbo.Products
WHERE ProductID < 60
GO
```

5. You can work with the values in the columns as they are returned. The following example performs a mathematical operation on the `Price` column. Columns that have been changed in this way will not have a name unless you provide one by using the `AS` keyword.

```
-- Returns ProductName and the Price including a 7% tax
-- Provides the name CustomerPays for the calculated column
SELECT ProductName, Price * 1.07 AS CustomerPays
FROM dbo.Products
GO
```

Functions That Are Useful in a SELECT Statement

For information about some functions that you can use to work with data in SELECT statements, see the following topics:

String Functions (Transact-SQL)	Date and Time Data Types and Functions (Transact-SQL)
Mathematical Functions (Transact-SQL)	Text and Image Functions (Transact-SQL)

Next Task in Lesson





[Summary: Creating Database Objects](#)

See Also

[SELECT \(Transact-SQL\)](#)

Lesson 1-5 - Summary - Creating Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

In this tutorial you have created a database, created a table in the database, inserted data into the table, changed the data, and then read the data from the table. The syntax for the statements that were used is only the basic syntax and many syntax options were not covered in this tutorial. To learn more about these statements, read the complete syntax for the statements in SQL Server Books Online, and review the many examples that are provided in those topics.

Next Lesson





[Lesson 2: Configuring Permissions on Database Objects](#)

See Also

[CREATE DATABASE \(SQL Server Transact-SQL\)](#)

Lesson 2: Configuring Permissions on Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Granting a user access to a database involves three steps. First, you create a login. The login lets the user connect to the SQL Server Database Engine. Then you configure the login as a user in the specified database. And finally, you grant that user permission to database objects. This lesson shows you these three steps, and shows you how to create a view and a stored procedure as the object.

This lesson contains the following topics:





- [Creating a Login](#)
- [Granting Access to a Database](#)
- [Creating Views and Stored Procedures](#)
- [Granting Access to a Database Object](#)
- [Summary: Configuring Permissions on Database Objects](#)

Next Task in Lesson

[Creating a Login](#)

Lesson 2-1 - Creating a Login

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

To access the Database Engine, users require a login. The login can represent the user's identity as a Windows account or as a member of a Windows group, or the login can be a SQL Server login that exists only in SQL Server. Whenever possible you should use Windows Authentication.

By default, administrators on your computer have full access to SQL Server. For this lesson, we want to have a less privileged user; therefore, you will create a new local Windows Authentication account on your computer. To do this, you must be an administrator on your computer. Then you will grant that new user access to SQL Server.

To create a new Windows account

1. Click **Start**, click **Run**, in the **Open** box, type **%SystemRoot%\system32\compmgmt.msc /s**, and then click **OK** to open the Computer Management program.
2. Under **System Tools**, expand **Local Users and Groups**, right-click **Users**, and then click **New User**.
3. In the **User name** box type **Mary**.
4. In the **Password** and **Confirm password** box, type a strong password, and then click **Create** to create a new local Windows user.

To create a login

1. In a Query Editor window of SQL Server Management Studio, type and execute the following code replacing `computer_name` with the name of your computer. `FROM WINDOWS` indicates that Windows will authenticate the user. The optional `DEFAULT_DATABASE = [TestData]` argument connects `Mary` to the `TestData` database, unless her connection string indicates another database. This statement introduces the semicolon as an optional termination for a Transact-SQL statement.

```
CREATE LOGIN [computer_name\Mary]
FROM WINDOWS
WITH DEFAULT_DATABASE = [TestData];
GO
```

This authorizes a user name `Mary`, authenticated by your computer, to access this instance of SQL Server. If there is more than one instance of SQL Server on the computer, you must create the login on each instance that `Mary` must access.

NOTE

Because `Mary` is not a domain account, this user name can only be authenticated on this computer.

Next Task in Lesson

[Granting Access to a Database](#)





See Also

CREATE LOGIN (Transact-SQL)

Choose an Authentication Mode

Lesson 2-2 - Granting Access to a Database

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Mary now has access to this instance of SQL Server, but does not have permission to access the databases. She does not even have access to her default database **TestData** until you authorize her as a database user.

To grant Mary access, switch to the **TestData** database, and then use the CREATE USER statement to map her login to a user named Mary.

To create a user in a database

1. Type and execute the following statements (replacing `computer_name` with the name of your computer) to grant `Mary` access to the `TestData` database.

```
USE [TestData];
GO

CREATE USER [Mary] FOR LOGIN [computer_name\Mary];
GO
```





Now, Mary has access to both SQL Server and the `TestData` database.

Next Task in Lesson

[Creating Views and Stored Procedures](#)

Lesson 2-3 - Creating Views and Stored Procedures

5/3/2018 • 2 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Now that Mary can access the **TestData** database, you may want to create some database objects, such as a view and a stored procedure, and then grant Mary access to them. A view is a stored SELECT statement, and a stored procedure is one or more Transact-SQL statements that execute as a batch.

Views are queried like tables and do not accept parameters. Stored procedures are more complex than views. Stored procedures can have both input and output parameters and can contain statements to control the flow of the code, such as IF and WHILE statements. It is good programming practice to use stored procedures for all repetitive actions in the database.

For this example, you will use CREATE VIEW to create a view that selects only two of the columns in the **Products** table. Then, you will use CREATE PROCEDURE to create a stored procedure that accepts a price parameter and returns only those products that cost less than the specified parameter value.

To create a view

1. Execute the following statement to create a very simple view that executes a select statement, and returns the names and prices of our products to the user.

```
CREATE VIEW vw_Names
AS
SELECT ProductName, Price FROM Products;
GO
```

Test the view

1. Views are treated just like tables. Use a `SELECT` statement to access a view.

```
SELECT * FROM vw_Names;
GO
```

To create a stored procedure

1. The following statement creates a stored procedure name `pr_Names`, accepts an input parameter named `@VarPrice` of data type `money`. The stored procedure prints the statement `Products less than` concatenated with the input parameter that is changed from the `money` data type into a `varchar(10)` character data type. Then, the procedure executes a `SELECT` statement on the view, passing the input parameter as part of the `WHERE` clause. This returns all products that cost less than the input parameter value.


```
CREATE PROCEDURE pr_Names @VarPrice money
AS
BEGIN
    -- The print statement returns text to the user
    PRINT 'Products less than ' + CAST(@VarPrice AS varchar(10));
    -- A second statement starts here
    SELECT ProductName, Price FROM vw_Names
        WHERE Price < @varPrice;
END
GO
```

Test the stored procedure

1. To test the stored procedure, type and execute the following statement. The procedure should return the names of the two products entered into the `Products` table in Lesson 1 with a price that is less than `10.00`.

```
EXECUTE pr_Names 10.00;
GO
```

Next Task in Lesson

[Granting Access to a Database Object](#)





See Also

[CREATE VIEW \(Transact-SQL\)](#)

[CREATE PROCEDURE \(Transact-SQL\)](#)

Lesson 2-4 - Granting Access to a Database Object

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

As an administrator, you can execute the SELECT from the **Products** table and the **vw_Names** view, and execute the **pr_Names** procedure; however, Mary cannot. To grant Mary the necessary permissions, use the GRANT statement.

Procedure Title

1. Execute the following statement to give `Mary` the `EXECUTE` permission for the `pr_Names` stored procedure.

```
GRANT EXECUTE ON pr_Names TO Mary;  
GO
```

In this scenario, Mary can only access the **Products** table by using the stored procedure. If you want Mary to be able to execute a SELECT statement against the view, then you must also execute

```
GRANT SELECT ON vw_Names TO Mary
```

 . To remove access to database objects, use the REVOKE statement.

NOTE

If the table, the view, and the stored procedure are not owned by the same schema, granting permissions becomes more complex.

About GRANT

You must have EXECUTE permission to execute a stored procedure. You must have SELECT, INSERT, UPDATE, and DELETE permissions to access and change data. The GRANT statement is also used for other permissions, such as permission to create tables.

Next Task in Lesson

[Summary: Configuring Permissions on Database Objects](#)





See Also

[GRANT \(Transact-SQL\)](#)

[REVOKE \(Transact-SQL\)](#)

Lesson 2-5 - Summary - Configuring Permissions on Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

Logins give users permissions to connect to SQL Server. Users are logins that can access a specific database. Use the GRANT statement to give users permission to read and to access and change the data.





A view is a single SELECT statement and looks like a table to the user. A stored procedure is one or more Transact-SQL statements that execute as a batch.

Next Lesson in Tutorial

[Lesson 3: Deleting Database Objects](#)

Lesson 3: Deleting Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

This short lesson removes the objects that you created in Lesson 1 and Lesson 2, and then drops the database.

This lesson contains one topic:





- [Deleting Database Objects](#)

Next Task in Lesson

[Deleting Database Objects](#)

Lesson 3-1 - Deleting Database Objects

5/3/2018 • 1 min to read • [Edit Online](#)

THIS TOPIC APPLIES TO:  SQL Server (starting with 2008)  Azure SQL Database  Azure SQL Data Warehouse  Parallel Data Warehouse

To remove all traces of this tutorial, you could just delete the database. However, in this topic, you will go through the steps to reverse every action you took doing the tutorial.

Removing permissions and objects

1. Before you delete objects, make sure you are in the correct database:

```
USE TestData;  
GO
```

2. Use the `REVOKE` statement to remove execute permission for `Mary` on the stored procedure:

```
REVOKE EXECUTE ON pr_Names FROM Mary;  
GO
```

3. Use the `DROP` statement to remove permission for `Mary` to access the `TestData` database:

```
DROP USER Mary;  
GO
```

4. Use the `DROP` statement to remove permission for `Mary` to access this instance of SQL Server 2005:

```
DROP LOGIN [<computer_name>\Mary];  
GO
```

5. Use the `DROP` statement to remove the store procedure `pr_Names`:

```
DROP PROC pr_Names;  
GO
```

6. Use the `DROP` statement to remove the view `vw_Names`:

```
DROP View vw_Names;  
GO
```

7. Use the `DELETE` statement to remove all rows from the `Products` table:

```
DELETE FROM Products;  
GO
```

8. Use the `DROP` statement to remove the `Products` table:

```
DROP Table Products;  
GO
```

9. You cannot remove the `TestData` database while you are in the database; therefore, first switch context to another database, and then use the `DROP` statement to remove the `TestData` database:

```
USE MASTER;  
GO  
DROP DATABASE TestData;  
GO
```

This concludes the Writing Transact-SQL Statements tutorial. Remember, this tutorial is a brief overview and it does not describe all the options to the statements that are used. Designing and creating an efficient database structure and configuring secure access to the data requires a more complex database than that shown in this tutorial.

Return to SQL Server Tools Portal

[Tutorial: Writing Transact-SQL Statements](#)

See Also

[REVOKE \(Transact-SQL\)](#)

[DROP USER \(Transact-SQL\)](#)

[DROP LOGIN \(Transact-SQL\)](#)

[DROP PROCEDURE \(Transact-SQL\)](#)

[DROP VIEW \(Transact-SQL\)](#)

[DELETE \(Transact-SQL\)](#)

[DROP TABLE \(Transact-SQL\)](#)

[DROP DATABASE \(Transact-SQL\)](#)